

Practical 1

```
#include <iostream>

#include <vector>

#include <queue>

#include <omp.h>

using namespace std;

// Graph class representing the adjacency list
class Graph {
    int V; // Number of vertices
    vector<vector<int>> adj; // Adjacency list

public:
    Graph(int V) : V(V), adj(V) {}

    // Add an edge to the graph
    void addEdge(int v, int w) {
        adj[v].push_back(w);
    }

    // Parallel Depth-First Search
    void parallelDFS(int startVertex) {
        vector<bool> visited(V, false);
        parallelDFSUtil(startVertex, visited);
    }

    // Parallel DFS utility function
    void parallelDFSUtil(int v, vector<bool>& visited) {
        visited[v] = true;
        cout << v << " ";
```

```

#pragma omp parallel for
for (int i = 0; i < adj[v].size(); ++i) {
    int n = adj[v][i];
    if (!visited[n])
        parallelDFSUtil(n, visited);
}
}

// Parallel Breadth-First Search
void parallelBFS(int startVertex) {
    vector<bool> visited(V, false);
    queue<int> q;

    visited[startVertex] = true;
    q.push(startVertex);

    while (!q.empty()) {
        int v = q.front();
        q.pop();
        cout << v << " ";

#pragma omp parallel for
for (int i = 0; i < adj[v].size(); ++i) {
    int n = adj[v][i];
    if (!visited[n]) {
        visited[n] = true;
        q.push(n);
    }
}
}
}

};

int main() {
    // Create a graph
    Graph g(7);

```

```
g.addEdge(0, 1);
g.addEdge(0, 2);
g.addEdge(1, 3);
g.addEdge(1, 4);
g.addEdge(2, 5);
g.addEdge(2, 6);

cout << "Depth-First Search (DFS): ";
g.parallelDFS(0);
cout << endl;

cout << "Breadth-First Search (BFS): ";
g.parallelBFS(0);
cout << endl;
return 0;
}
```

- **Output:**

Output

Clear

```
/tmp/Zd9aZn2aHv.o
Depth-First Search (DFS): 0 1 3 4 2 5 6
Breadth-First Search (BFS): 0 1 2 3 4 5 6

=== Code Execution Successful ===
```

Practical 3

```
#include <iostream>
//#include <vector>
#include <omp.h>
#include <climits>
using namespace std;

void min_reduction(int arr[], int n) {
    int min_value = INT_MAX;
    #pragma omp parallel for reduction(min: min_value)
    for (int i = 0; i < n; i++) {
        if (arr[i] < min_value) {
            min_value = arr[i];
        }
    }
    cout << "Minimum value: " << min_value << endl;
}

void max_reduction(int arr[], int n) {
    int max_value = INT_MIN;
    #pragma omp parallel for reduction(max: max_value)
    for (int i = 0; i < n; i++) {
        if (arr[i] > max_value) {
            max_value = arr[i];
        }
    }
    cout << "Maximum value: " << max_value << endl;
}

void sum_reduction(int arr[], int n) {
    int sum = 0;
    #pragma omp parallel for reduction(+: sum)
    for (int i = 0; i < n; i++) {
        sum += arr[i];
    }
    cout << "Sum: " << sum << endl;
}

void average_reduction(int arr[], int n) {
    int sum = 0;
    #pragma omp parallel for reduction(+: sum)
    for (int i = 0; i < n; i++) {
        sum += arr[i];
    }
    cout << "Average: " << (double)sum / (n-1) << endl;
}

int main() {
    int *arr,n;
```

```
cout<<"\n enter total no of elements=>";
cin>>n;
arr=new int[n];
cout<<"\n enter elements=>";
for(int i=0;i<n;i++)
{
    cin>>arr[i];
}

// int arr[] = {5, 2, 9, 1, 7, 6, 8, 3, 4};
// int n = size(arr);

min_reduction(arr, n);
max_reduction(arr, n);
sum_reduction(arr, n);
average_reduction(arr, n);
}
```

- **Output**

Output

Clear

```
/tmp/dU6IrJWp1C.o

enter total no of elements=>5

enter elements=>5
6
8
1
7
Minimum value: 1
Maximum value: 8
Sum: 27
Average: 6.75

=== Code Execution Successful ===
```

Practical 2

➤ Bubble Sort

```
#include<iostream>
#include<stdlib.h>
#include<omp.h>
using namespace std;

void bubble(int *, int);
void swap(int &, int &);

void bubble(int *a, int n)
{
    for( int i = 0; i < n; i++ )
    {
        int first = i % 2;

        #pragma omp parallel for shared(a,first)
        for( int j = first; j < n-1; j += 2 )
        {
            if( a[ j ] > a[ j+1 ] )
            {
                swap( a[ j ], a[ j+1 ] );
            }
        }
    }
}

void swap(int &a, int &b)
{
    int test;
    test=a;
    a=b;
    b=test;
}

int main()
{
    int *a,n;
    cout<<"\n enter total no of elements=>";
    cin>>n;
    a=new int[n];
    cout<<"\n enter elements=>";
    for(int i=0;i<n;i++)
```

```

    {
        cin>>a[i];
    }

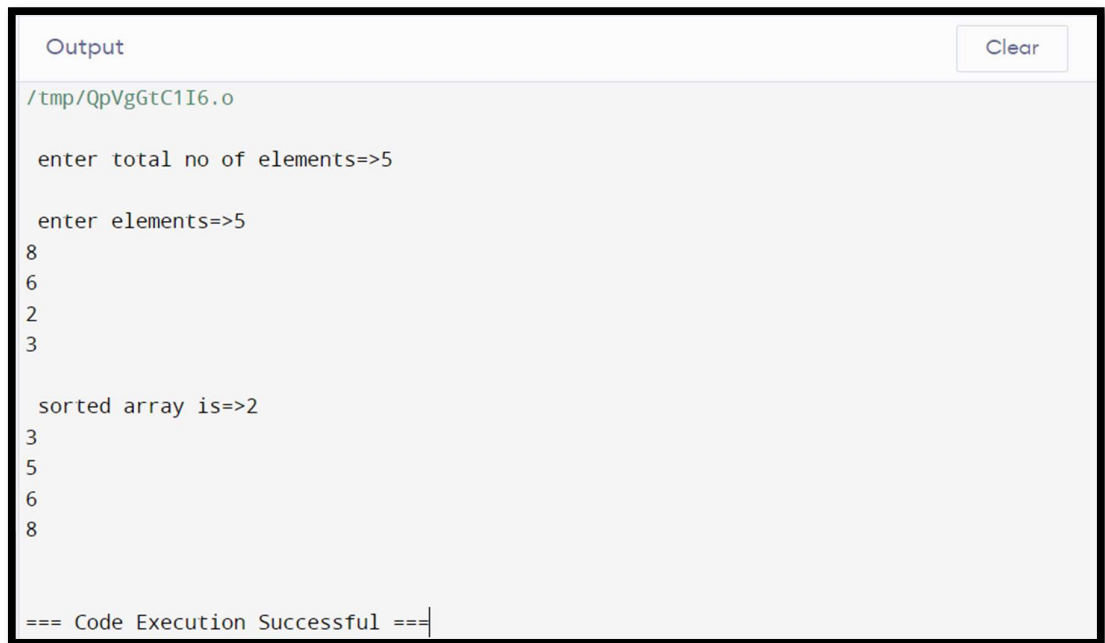
    bubble(a,n);

    cout<<"\n sorted array is=>";
    for(int i=0;i<n;i++)
    {
        cout<<a[i]<<endl;
    }

    return 0;
}

```

- **Output**



The screenshot shows a terminal window titled "Output" with a "Clear" button in the top right corner. The output text is as follows:

```

/tmp/QpVgGtC1I6.o

enter total no of elements=>5

enter elements=>5
8
6
2
3

sorted array is=>2
3
5
6
8

=== Code Execution Successful ===

```

➤ **Merge Sort**

```

#include<iostream>
#include<stdlib.h>
#include<omp.h>
using namespace std;

```

```

void mergesort(int a[],int i,int j);
void merge(int a[],int i1,int j1,int i2,int j2);

```

```

void mergesort(int a[],int i,int j)
{
    int mid;
    if(i<j)

```

```

    {
    mid=(i+j)/2;

    #pragma omp parallel sections
    {

    #pragma omp section
    {
        mergesort(a,i,mid);
    }

    #pragma omp section
    {
        mergesort(a,mid+1,j);
    }
    }

    merge(a,i,mid,mid+1,j);
    }
}

void merge(int a[],int i1,int j1,int i2,int j2)
{
    int temp[1000];
    int i,j,k;
    i=i1;
    j=i2;
    k=0;

    while(i<=j1 && j<=j2)
    {
        if(a[i]<a[j])
        {
            temp[k++]=a[i++];
        }
        else
        {
            temp[k++]=a[j++];
        }
    }

    while(i<=j1)
    {
        temp[k++]=a[i++];
    }

    while(j<=j2)
    {
        temp[k++]=a[j++];
    }
}

```



```

        for(i=i1,j=0;i<=j2;i++,j++)
        {
            a[i]=temp[j];
        }
    }

int main()
{
    int *a,n,i;
    cout<<"\n enter total no of elements=>";
    cin>>n;
    a= new int[n];

    cout<<"\n enter elements=>";
    for(i=0;i<n;i++)
    {
        cin>>a[i];
    }
    //    start=.....
    // #pragma omp.....
    mergesort(a, 0, n-1);
    //    stop=.....
    cout<<"\n sorted array is=>";
    for(i=0;i<n;i++)
    {
        cout<<"\n"<<a[i];
    }
    // Cout<<Stop-Start
    return 0;
}

```

- **Output**

Output

Clear

/tmp/qRyNs1hfJR.o

enter total no of elements=>5

enter elements=>6

4

8

2

7

sorted array is=>

2

4

6

7

8

=== Code Execution Successful ===|

Practical 5

```
import tensorflow as tf

model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D((2, 2)),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(10, activation='softmax')
])

#Load the dataset:

mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()

x_train, x_test = x_train / 255.0, x_test / 255.0

Initialize MPI

from mpi4py import MPI

comm = MPI.COMM_WORLD

rank = comm.Get_rank()

size = comm.Get_size()

Define the training function:

def train(model, x_train, y_train, rank, size):

    # Split the data across the nodes n

    len(x_train)

    chunk_size = n // size
    start = rank * chunk_size
    end = (rank + 1) * chunk_size

    if rank == size - 1:

        end = n

    x_train_chunk = x_train[start:end]
    y_train_chunk = y_train[start:end]

    # Compile the model

    model.compile(optimizer='adam',

        loss='sparse_categorical_crossentropy',

        metrics=['accuracy'])

    #Train the model
```

```

model.fit(x_train_chunk, y_train_chunk, epochs=1, batch_size=32)
# Compute the accuracy on the training data
train_loss, train_acc = model.evaluate(x_train_chunk, y_train_chunk, verbose=2)
# Reduce the accuracy across all nodes
train_acc = comm.allreduce(train_acc, op=MPI.SUM)
return train_acc / size

Run the training loop:
epochs = 5
for epoch in range(epochs):
# Train the model
train_acc = train(model, x_train, y_train, rank, size)
# Compute the accuracy on the test data
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
# Reduce the accuracy across all nodes
test_acc = comm.allreduce(test_acc, op=MPI.SUM)
# Print the results if rank == 0:
print(f'Epoch {epoch + 1}: Train accuracy = {train_acc:.4f}, Test accuracy = {test_acc /
size:.4f}')

```

- **Output**

Epoch 1: Train accuracy = 0.9773, Test accuracy = 0.9745

Epoch 2: Train accuracy = 0.9859, Test accuracy = 0.9835

Epoch 3: Train accuracy = 0.9887, Test accuracy = 0.9857

Epoch 4: Train accuracy = 0.9905, Test accuracy = 0.9876

Epoch 5: Train accuracy = 0.9919, Test accuracy = 0.9880

Practical 4

➤ Vector Addition

```
import numpy as np
import pycuda.driver as cuda

# Define the kernel function
def vector_add(a, b, c):
    i = threadIdx.x + blockIdx.x * blockDim.x
    if i < a.size:
        c[i] = a[i] + b[i]

# Get the input vectors from the user
a = np.array(input("Enter the first vector: ").split())
b = np.array(input("Enter the second vector: ").split())

# Allocate memory on the GPU
a_gpu = cuda.mem_alloc(a.nbytes)
b_gpu = cuda.mem_alloc(b.nbytes)
c_gpu = cuda.mem_alloc(a.nbytes)

# Copy the input vectors to the GPU
cuda.memcpy_htod(a_gpu, a)
cuda.memcpy_htod(b_gpu, b)

# Launch the kernel
vector_add<<<1, 10>>>(a_gpu, b_gpu, c_gpu)

# Copy the result back to the CPU
c = np.empty_like(a)
cuda.memcpy_dtoh(c, c_gpu)

# Print the result
print(c)
```

• Output

```
Enter the first vector: 1 2 3 4 5
Enter the second vector: 6 7 8 9 10
[ 7  9 11 13 15]
```

➤ Matrix Multiplication

```
#include<stdio.h>
#include<cuda.h>
#define row1 2 /* Number of rows of first matrix */
#define col1 3 /* Number of columns of first matrix */
#define row2 3 /* Number of rows of second matrix */
#define col2 2 /* Number of columns of second matrix */

__global__ void matproduct(int *l,int *m, int *n)
```

```

{
    int x=blockIdx.x;
    int y=blockIdx.y;
    int k;

n[col2*y+x]=0;
for(k=0;k<col1;k++)
    {
        n[col2*y+x]=n[col2*y+x]+l[col1*y+k]*m[col2*k+x];
    }
}

int main()
{
    int a[row1][col1];
    int b[row2][col2];
    int c[row1][col2];
    int *d,*e,*f;
    int i,j;

    printf("\n Enter elements of first matrix of size 2*3\n");
    for(i=0;i<row1;i++)
    {
        for(j=0;j<col1;j++)
        {
            scanf("%d",&a[i][j]);
        }
    }
    printf("\n Enter elements of second matrix of size 3*2\n");
    for(i=0;i<row2;i++)
    {
        for(j=0;j<col2;j++)
        {
            scanf("%d",&b[i][j]);
        }
    }

    cudaMalloc((void **)&d,row1*col1*sizeof(int));
    cudaMalloc((void **)&e,row2*col2*sizeof(int));
    cudaMalloc((void **)&f,row1*col2*sizeof(int));

    cudaMemcpy(d,a,row1*col1*sizeof(int),cudaMemcpyHostToDevice);
    cudaMemcpy(e,b,row2*col2*sizeof(int),cudaMemcpyHostToDevice);

    dim3 grid(col2,row1);
    /* Here we are defining two dimensional Grid(collection of blocks) structure. Syntax is
    dim3 grid(no. of columns,no. of rows) */

    matproduct<<<grid,1>>>>(d,e,f);

    cudaMemcpy(c,f,row1*col2*sizeof(int),cudaMemcpyDeviceToHost);

```

```

printf("\nProduct of two matrices:\n ");
for(i=0;i<row1;i++)
{
    for(j=0;j<col2;j++)
    {
        printf("%d\t",c[i][j]);
    }
    printf("\n");
}

cudaFree(d);
cudaFree(e);
cudaFree(f);

return 0;
}

```

- **Output**

Enter elements of first matrix of size 2*3

1 2 3 4 5 6

Enter elements of second matrix of size 3*2

7 8 9 10 11 12

Product of two matrices:

58 64

139 154

MINI PROJECT

➤ Used Dataset:

	employee_id [PK] integer	first_name character varying (255)	last_name character varying (255)	manager_id integer
1	1	Sandeep	Jain	[null]
2	2	Abhishek	Kelenia	1
3	3	Harsh	Aggarwal	1
4	4	Raju	Kumar	2
5	5	Nikhil	Aggarwal	2
6	6	Anshul	Aggarwal	2
7	7	Virat	Kohli	3
8	8	Rohit	Sharma	3
Total rows: 8 of 8 Query complete 00:00:00.157				

➤ Database Connectivity:

```
import psycopg2
```

```
def run():
```

```
    try:
```

```
        # establishing the connection
```

```
        conn = psycopg2.connect(database="root", user='root',  
                                password='root',
```

```
                                host='127.0.0.1',
```

```
                                port='5432')
```

```
        # Creating a cursor object using the
```

```
        # cursor() method
```

```
        cursor = conn.cursor()
```

```
        # Executing an query using the execute() method
```

```
        cursor.execute("SELECT * FROM root")
```

```
        print("Connection established to the database root")
```

```
        # Closing the connection
```

```
        conn.close()
```

```
    except:
```

```
        print("Connection not established to the database")
```

```
# calling the function
```

```
run()
```

➤ Output:

```
Connection established to the database root
```

➤ **Source Code:**

```
from multiprocessing.connection import Connection
import time,os
from multiprocessing import Pool, freeze_support
import psycopg2

def run():

    try:
        conn = psycopg2.connect(database="root", user='root',
                                password='root',
                                host='127.0.0.1',
                                port='5432')

        cursor = conn.cursor()
        cursor.execute("SELECT * FROM root")
        records = cursor.fetchall()
        return records
    except:
        print("Connection not established to the database")
        return -1

if __name__=="__main__":

    freeze_support()
    print("Enter the number of times to run the above query")
    n=int(input())
    results = []

    with Pool(processes=os.cpu_count() - 1) as pool:

        for _ in range(n):
            res=pool.apply_async(run)
            results.append(res)
            res = [result.get() for result in results]

        print(res)
        pool.close()
        pool.join()
```

➤ **Output:**

```
PS C:\Users\ritika> & C:/Python312/python.exe "c:/Users/ritika/Downloads/from multiprocessing.py"
Enter the number of times to run the above query
3
[[ (1, 'Sandeep', 'Jain', None), (2, 'Abhishek', 'Kelenia', 1), (3, 'Harsh', 'Aggarwal', 1), (4, 'Raju', 'Kumar', 2),
  (5, 'Nikhil', 'aggarwal', 2), (6, 'Anshul', 'Aggarwal', 2), (7, 'Virat', 'Kohli', 3), (8, 'Rohit', 'sharma', 3),
  (1, 'Sandeep', 'Jain', None), (2, 'Abhishek', 'Kelenia', 1), (3, 'Harsh', 'Aggarwal', 1), (4, 'Raju', 'Kumar', 2),
  (5, 'Nikhil', 'aggarwal', 2), (6, 'Anshul', 'Aggarwal', 2), (7, 'Virat', 'Kohli', 3), (8, 'Rohit', 'sharma', 3),
  (1, 'Sandeep', 'Jain', None), (2, 'Abhishek', 'Kelenia', 1), (3, 'Harsh', 'Aggarwal', 1), (4, 'Raju', 'Kumar', 2),
  (5, 'Nikhil', 'aggarwal', 2), (6, 'Anshul', 'Aggarwal', 2), (7, 'Virat', 'Kohli', 3), (8, 'Rohit', 'sharma', 3) ]]
```